VICTOR IBHAFIDON

Finance Systems & Process Business Analyst

London, United Kingdom | victoribhafidon@outlook.com | 07425 871413 | linkedin.com/in/victoribhafidon | github.com/Web8080

PROFESSIONAL SUMMARY

Engineer with MSc Data Science and hands-on analytical delivery. Designed PostgreSQL data architecture and analytics pipelines for an MVP platform, and built behavioural anomaly detection reaching 0.981 ROC-AUC for an AI governance platform. Comfortable turning raw data into structured analysis, metrics and recommendations for non-technical stakeholders. Combines statistical skill with the ability to own work from data collection through to insight and action.

TECHNICAL SKILLS

AI & Machine Learning: LLMs, Generative AI, AI Agents, Multi-Agent Systems, LangGraph, MCP, JSON-RPC, Tool Calling, RAG, GraphRAG, Prompt Engineering, AI Evaluation, PyTorch, Transformers, Scikit-learn, Computer Vision, NLP

Backend & Data: FastAPI, Django, REST APIs, WebSockets, Microservices, PostgreSQL, MySQL, MongoDB, Redis, Neo4j, pgvector, Vector Search, Data Pipelines, ETL

MLOps & Model Operations: MLflow, Model Training, Model Evaluation, Model Monitoring, Drift Detection, Experiment Tracking, Model Deployment

Programming: Python, TypeScript, JavaScript, SQL, Bash

Security & AI Governance: AI Security, AI Governance, Authentication, Authorization, RBAC, API Security, OWASP Top 10, Prompt Injection Mitigation, AI Safety, Responsible AI, GDPR

Cloud & DevOps: AWS (EC2, S3, Lambda, IAM, EKS), Microsoft Azure, Azure OpenAI, Docker, Kubernetes, GitHub Actions, CI/CD, Linux, Nginx, OpenTelemetry, Prometheus, Grafana

Analysis & BI: Data Analysis, SQL, Pandas, Excel, KPI Reporting, Dashboards, Statistical Analysis, Anomaly Detection, Feature Engineering, Data Pipelines, ETL

Business & Communication: Stakeholder Reporting, Insight Communication, Requirements Translation

PROFESSIONAL EXPERIENCE

Software Engineer (AI/ML), Bastion Shield Technologies

London, United Kingdom | April 2026 – Present

- At Bastion Shield Technologies I designed the data and analytics layer of Bulwark, an AI governance platform, collecting, validating and structuring agent activity, permissions, policy violations and risk signals. I applied feature engineering, statistical analysis and machine learning to detect behavioural patterns and anomalies (0.981 ROC-AUC), and translated the results into clear, actionable insights for security and product teams.
- Developed behavioural anomaly detection combining statistical methods with usage forecasting to identify unusual agent activity. The system achieved a 0.981 ROC-AUC against an Isolation Forest baseline and was integrated into the TypeScript runtime for real-time scoring.
- Engineered the backend in Python, TypeScript, and PostgreSQL, implementing RBAC, server-side tenant isolation, secure authentication, API-key management, and an SDK for external teams to register agents and stream policy decisions. Designed tamper-evident audit logging with offline verification, strengthened data privacy controls, and remediated four production vulnerabilities, including a session-token forgery path.
- Owned the cloud infrastructure using Terraform-managed AWS environments across development, staging, and production. Containerised the platform with Docker and delivered it through GitHub Actions CI/CD, while working on distributed system patterns including queues, event streaming, caching, and horizontal scaling to improve performance and latency.

Full Stack Software Engineer, Basketball Nxtion

London, United Kingdom | February 2026 – April 2026

- Developed the MVP of Plxyground, a mobile application connecting basketball players, creators, and sports communities, contributing across product design, application architecture, backend development, APIs, authentication, and user workflows.
- Designed the PostgreSQL data architecture and analytics pipelines for the MVP, structuring user, activity, engagement, and platform data to support data ingestion, transformation, analysis, and business intelligence. Built data workflows enabling insights into platform performance, user behaviour, engagement, and product adoption.
- Designed and integrated an AI architecture connecting LLM/ML services with PostgreSQL data, backend APIs, and application business logic. Built workflows for personalised recommendations, intelligent user matching, and content discovery, using user and platform data to generate relevant results, improve engagement, and automate discovery workflows.
- Deployed and maintained the MVP infrastructure on AWS, implementing Docker, CI/CD, testing, rollbacks, monitoring, and third-party integrations while collaborating with frontend engineers and stakeholders throughout the software development lifecycle from requirements and development through testing, deployment, and iteration.

Full Stack Python Engineer Intern, CyberAI Technologies

London, United Kingdom | September 2025 – November 2025

- Collaborated with senior AI engineers to develop Surfyn AI, a conversational AI customer support platform enabling businesses to deploy knowledge-based AI assistants grounded in websites and documents to automate customer interactions and provide relevant, context-aware responses.
- Contributed to the development of RAG and conversational AI workflows using Python and Django, working on document ingestion, chunking, embeddings, vector-based retrieval, prompt construction, conversation context, REST APIs, and LLM integrations while learning how knowledge-based AI systems are designed and evaluated.
- Developed backend functionality and AI workflow improvements, working on API development, debugging, testing, performance optimisation, knowledge retrieval, and integration between AI services and the Django application.
- Worked with the AI Security Engineering team on AI evaluation, testing, safety, and deployment workflows, helping investigate retrieval quality, response relevance, hallucination risks, prompt injection, sensitive data handling, logging, and application reliability across production AI systems.

SELECTED PROJECTS

Portivo Estate Suite: Took ownership of the system design and development of an AI powered property operations platform for UK letting agents and BTR operators, covering lettings, sales, maintenance and compliance in one system. Built Portivo end to end with a React, TypeScript and Tailwind frontend coupled with a Python and FastAPI backend using SQLAlchemy and PostgreSQL for the multi-tenant data model. Implemented Alembic schema migrations against a live customer database, Redis and Celery for caching and scheduled background jobs, tenant isolation at the query layer, and role-based access control at the endpoint level. The AI layer uses a LangGraph multi-agent workflow on Claude models, with an orchestrator delegating typed tasks to specialist agents for document retrieval, lease abstraction, arrears detection, compliance tracking, maintenance triage and outbound sales prospecting. Built an OCR, chunking and embedding pipeline using pgvector with per-tenant vector scoping and source-referenced answers, while agents access CRM, email and document capabilities through MCP servers with scoped credentials rather than bespoke integrations. Evaluated against labelled datasets for groundedness and extraction accuracy, traced with Langfuse, and guarded using schema-validated outputs and confidence thresholds. Deployed to the cloud using Docker and GitHub Actions CI/CD, with an audit trail covering every agent action. Live at portivoestatesuite.com

Medical Imaging AI: Developed a supervised computer vision system for classifying 231,444 medical images, designing the data preparation and augmentation pipeline to improve the quality and consistency of training data and reduce sensitivity to variations in the input images. Iteratively trained and evaluated PyTorch CNN models, analysing classification metrics and failure patterns to guide model and preprocessing changes, achieving up to 92% classification accuracy. Converted the trained model into a deployable inference service using Docker and REST APIs, separating model training from inference so models could be versioned and deployed independently. Added health checks, audit trails, and runtime monitoring to make model serving more observable and reliable beyond the training environment. GitHub: github.com/Web8080/Medical_Imaging_AI_API_Research_Paper_and_web_app

EDUCATION

MSc Data Science, Nottingham Trent University, United Kingdom (2023 – 2024)

BSc Computer Science, Ambrose Alli University, Nigeria (2016 – 2021)