

VICTOR IBHAFIDON

Marketing Science Analyst

London, United Kingdom | victoribhafidon@outlook.com | 07425 871413 | linkedin.com/in/victoribhafidon | github.com/Web8080

PROFESSIONAL SUMMARY

Engineer with hands-on analytical delivery. Designed PostgreSQL data architecture and analytics pipelines for an MVP platform, and built behavioural anomaly detection reaching 0.981 ROC-AUC for an AI governance platform. Comfortable turning raw data into structured analysis, metrics and recommendations for non-technical stakeholders. Combines statistical skill with the ability to own work from data collection through to insight and action. I have worked closely with statistical and machine-learning models in production, integrating trained models into live runtimes for real-time scoring and evaluating them against baselines rather than building them in isolation.

TECHNICAL SKILLS

Backend & Data: Data Pipelines, FastAPI, Django, REST APIs, WebSockets, Microservices, PostgreSQL, MySQL, MongoDB, Redis, Neo4j, pgvector, Vector Search, ETL

Cloud & DevOps: IAM, AWS (EC2, S3, Lambda, EKS), Microsoft Azure, Azure OpenAI, Docker, Kubernetes, GitHub Actions, CI/CD, Linux, Nginx, OpenTelemetry, Prometheus, Grafana

AI & Machine Learning: LLMs, Generative AI, AI Agents, Multi-Agent Systems, LangGraph, MCP, JSON-RPC, Tool Calling, RAG, GraphRAG, Prompt Engineering, AI Evaluation, PyTorch, Transformers, Scikit-learn, Computer Vision, NLP

MLOps & Model Operations: MLflow, Model Training, Model Evaluation, Model Monitoring, Drift Detection, Experiment Tracking, Model Deployment

Programming: Python, TypeScript, JavaScript, SQL, Bash

Security & AI Governance: AI Security, AI Governance, Authentication, Authorization, RBAC, API Security, OWASP Top 10, Prompt Injection Mitigation, AI Safety, Responsible AI, GDPR

Analysis & BI: Data Analysis, SQL, Pandas, Excel, KPI Reporting, Dashboards, Statistical Analysis, Anomaly Detection, Feature Engineering, Data Pipelines, ETL

Business & Communication: Stakeholder Reporting, Insight Communication, Requirements Translation

PROFESSIONAL EXPERIENCE

Software Engineer (AI/ML), Bastion Shield Technologies

London, United Kingdom | April 2026 – Present

- At Bastion Shield Technologies I designed the data and analytics layer of Bulwark, an AI governance platform, collecting, validating and structuring agent activity, permissions, policy violations and risk signals. I applied feature engineering, statistical analysis and machine learning to detect behavioural patterns and anomalies (0.981 ROC-AUC), and translated the results into clear, actionable insights for security and product teams.
- Developed behavioural anomaly detection combining statistical methods with usage forecasting to identify unusual agent activity. The system achieved a 0.981 ROC-AUC against an Isolation Forest baseline and was integrated into the TypeScript runtime for real-time scoring.
- Engineered the backend in Python, TypeScript, and PostgreSQL, implementing RBAC, server-side tenant isolation, secure authentication, API-key management, and an SDK for external teams to register agents and stream policy decisions. Designed tamper-evident audit logging with offline verification, strengthened data privacy controls, and remediated four production vulnerabilities, including a session-token forgery path.
- Owned the cloud infrastructure using Terraform-managed AWS environments across development, staging, and production. Containerised the platform with Docker and delivered it through GitHub Actions CI/CD, while working on distributed system patterns including queues, event streaming, caching, and horizontal scaling to improve performance and latency.

Full Stack Software Engineer, Basketball Nxtion

London, United Kingdom | February 2026 – April 2026

- Developed the MVP of Pxyground, a mobile application connecting basketball players, creators, and sports communities, contributing across product design, application architecture, backend development, APIs, authentication, and user workflows.
- Designed the PostgreSQL data architecture and analytics pipelines for the MVP, structuring user, activity, engagement, and platform data to support data ingestion, transformation, analysis, and business intelligence. Built data workflows enabling insights into platform performance, user behaviour, engagement, and product adoption.
- Designed and integrated an AI architecture connecting LLM/ML services with PostgreSQL data, backend APIs, and application business logic. Built workflows for personalised recommendations, intelligent user matching, and content discovery, using user and platform data to generate relevant results, improve engagement, and automate discovery workflows.

- Deployed and maintained the MVP infrastructure on AWS, implementing Docker, CI/CD, testing, rollbacks, monitoring, and third-party integrations while collaborating with frontend engineers and stakeholders throughout the software development lifecycle from requirements and development through testing, deployment, and iteration.

Full Stack Python Engineer Intern, CyberAI Technologies

London, United Kingdom | September 2025 – November 2025

- Collaborated with senior AI engineers to develop Surfyn AI, a conversational AI customer support platform enabling businesses to deploy knowledge-based AI assistants grounded in websites and documents to automate customer interactions and provide relevant, context-aware responses.
- Contributed to the development of RAG and conversational AI workflows using Python and Django, working on document ingestion, chunking, embeddings, vector-based retrieval, prompt construction, conversation context, REST APIs, and LLM integrations while learning how knowledge-based AI systems are designed and evaluated.
- Developed backend functionality and AI workflow improvements, working on API development, debugging, testing, performance optimisation, knowledge retrieval, and integration between AI services and the Django application.
- Worked with the AI Security Engineering team on AI evaluation, testing, safety, and deployment workflows, helping investigate retrieval quality, response relevance, hallucination risks, prompt injection, sensitive data handling, logging, and application reliability across production AI systems.

SELECTED PROJECTS

Portivo Estate Suite (AI Powered PropTech Operating System): Approached Portivo from a consulting and product-engineering perspective, working with estate managers to understand how AI could be integrated into existing property operations to reduce administrative workload and improve lettings, sales, maintenance, compliance and reporting. Worked directly with users through discovery, workflow design, testing and implementation, translating feedback into product requirements, documenting APIs and operational processes, validating end-to-end workflows and supporting customer configuration and adoption. Designed the system architecture and took ownership of end-to-end development across product analysis, UX/UI, frontend, backend, AI/ML, data engineering, cloud, security, QA, AI governance and integrations. Built the multi-tenant application with React, TypeScript, Next.js and Tailwind, developing dashboards, property management, CRM, reporting and administrative workflows, backed by Python, FastAPI and SQLAlchemy for business logic and APIs. Designed the PostgreSQL data model with tenant-scoped queries and Alembic migrations against a live customer database, using Redis and Celery for caching, asynchronous processing and scheduled jobs. Built the AI architecture using LangGraph, Anthropic Models, RAG, OCR, chunking, embeddings and pgvector, with an orchestrator delegating typed tasks to specialist agents for document retrieval, lease abstraction, arrears detection, compliance tracking, maintenance triage and outbound sales prospecting. Implemented MCP servers and scoped credentials to give agents controlled access to CRM, email and document tools rather than creating bespoke integrations. Evaluated AI performance using labelled datasets for groundedness and extraction accuracy, with Langfuse for tracing, schema-validated outputs, confidence thresholds and audit trails for agent actions. Applied RBAC, tenant isolation, scoped permissions, security controls and AI governance to protect customer data and control agent behaviour, including prompt-injection and data-leakage risks. Containerised and deployed the platform to the cloud using Docker and GitHub Actions CI/CD, while designing API integrations, technical documentation and implementation workflows to support enterprise adoption. **Live in production at:** portivoestatesuite.com

E-Commerce Sales Analytics Platform: An end-to-end analytics platform built to turn two years of raw e-commerce sales data into decisions the business could act on. Working from CSV exports covering 45,958 transactions, 50,000 customers and \$33M of revenue, it answers the questions a commercial team actually asks: who is about to churn, which segments are worth retaining, and where revenue is concentrated. Technically it spans the full analytics lifecycle: data engineering and cleaning in Python and SQL over PostgreSQL, exploratory and statistical analysis, RFM segmentation, cohort and time-series analysis, customer lifetime value modelling and churn prediction, surfaced through interactive dashboards and executive reporting. GitHub: github.com/Web8080/Ecommerce_Sales_Analytics

EDUCATION

MSc Data Science, Nottingham Trent University, United Kingdom (2023 – 2024)

BSc Computer Science, Ambrose Alli University, Nigeria (2016 – 2021)